

```
# Topic: Implement KNN in Water Quality Dataset
# Collaborated by : Aadith Joseph Mathew and Devika S Vinod
# Roll No: 23122101, 23122113
# Date: 16th May 2024
# Submission : 17th May 2024
```

Dataset Description:

- Rows: 7,999
- Columns: 21
- Features: Chemical concentrations: Includes various chemicals like aluminium, ammonia, arsenic, barium, cadmium, chloramine, chromium, copper, flouride, lead, nitrates, nitrites, mercury, perchlorate, radium, selenium, silver, and uranium. Biological parameters: bacteria and viruses, indicating the presence of biological contaminants. Safety Indicator: is_safe (1 for safe, 0 for not safe), indicating whether the water quality is considered safe for consumption based on the measured parameters.

References:

- <https://www.kaggle.com/code/sabrinajeannin/wine-pca-hierarchical-clustering>
- <https://www.kaggle.com/code/elisthefox/ultimate-guide-to-k-nearest-neighbors-k-nn#4.-Results-of-K-NN-implementation-after-preparation-of-data>

Tools and Libraries:

- Pandas: For data manipulation and analysis.
- Matplotlib: For creating visualizations such as plots.
- Seaborn: For enhancing the visual aesthetics of plots.
- Scikit-learn: For building and evaluating machine learning models.

Sections:

- Importing Libraries
 - Loading and Viewing the Dataset
 - EDA
 - Visualization
 - Model Building
 - Model Improvement
 - Hyperparameter Tuning
-

Importing Libraries

```
# Importing required libraries
import pandas as pd
import numpy as np
import matplotlib.pyplot as plt
import seaborn as sns
```

```

from sklearn.model_selection import train_test_split
from sklearn.neighbors import KNeighborsClassifier
from sklearn.metrics import accuracy_score, confusion_matrix,
classification_report

```

Loading and Viewing the dataset

```

# Loading dataset
df = pd.read_csv("waterQuality1.csv")

# First 5 rows
df.head()

```

	aluminium	ammonia	arsenic	barium	cadmium	chloramine	chromium
0	1.65	9.08	0.04	2.85	0.007	0.35	0.83
1	2.32	21.16	0.01	3.31	0.002	5.28	0.68
2	1.01	14.02	0.04	0.58	0.008	4.24	0.53
3	1.36	11.33	0.04	2.96	0.001	7.23	0.03
4	0.92	24.33	0.03	0.20	0.006	2.67	0.69

	flouride perchlorate	bacteria	...	lead	nitrates	nitrites	mercury
0	0.05	0.20	...	0.054	16.08	1.13	0.007
1	0.90	0.65	...	0.100	2.01	1.93	0.003
2	0.99	0.05	...	0.078	14.16	1.11	0.006
3	1.08	0.71	...	0.016	1.41	1.29	0.004
4	0.61	0.13	...	0.117	6.74	1.11	0.003

	radium	selenium	silver	uranium	is_safe
0	6.78	0.08	0.34	0.02	1
1	3.21	0.08	0.27	0.05	1
2	7.07	0.07	0.44	0.01	0
3	1.72	0.02	0.45	0.05	1
4	2.41	0.02	0.06	0.02	1

[5 rows x 21 columns]

```

# Last 5 rows
df.tail()

```

	aluminium	ammonia	arsenic	barium	cadmium	chloramine
chromium \						
7991	0.05	7.78	0.00	1.95	0.04	0.10
0.03						
7992	0.05	24.22	0.02	0.59	0.01	0.45
0.02						
7993	0.09	6.85	0.00	0.61	0.03	0.05
0.05						
7994	0.01	10.00	0.01	2.00	0.00	2.00
0.00						
7995	0.04	6.85	0.01	0.70	0.03	0.05
0.01						

	copper	flouride	bacteria	...	lead	nitrates	nitrites
mercury \							
7991	0.03	1.37	0.0	...	0.197	14.29	1.0
0.005							
7992	0.02	1.48	0.0	...	0.031	10.27	1.0
0.001							
7993	0.02	0.91	0.0	...	0.182	15.92	1.0
0.000							
7994	0.09	0.00	0.0	...	0.000	0.00	0.0
0.000							
7995	0.03	1.00	0.0	...	0.182	15.92	1.0
0.000							

	perchlorate	radium	selenium	silver	uranium	is_safe
7991	3.57	2.13	0.09	0.06	0.03	1
7992	1.48	1.11	0.09	0.10	0.08	1
7993	1.35	4.84	0.00	0.04	0.05	1
7994	0.00	0.00	0.00	0.00	0.00	1
7995	1.35	4.84	0.00	0.04	0.05	1

[5 rows x 21 columns]

Getting the dimension of the dataset

df.shape

(7996, 21)

EDA

Getting column names

df.columns

```
Index(['aluminium', 'ammonia', 'arsenic', 'barium', 'cadmium',
      'chloramine',
      'chromium', 'copper', 'flouride', 'bacteria', 'viruses',
      'lead',
      'nitrates', 'nitrites', 'mercury', 'perchlorate', 'radium',
```

```
'selenium',
      'silver', 'uranium', 'is_safe'],
      dtype='object')
```

```
# Getting column information
```

```
df.info()
```

```
<class 'pandas.core.frame.DataFrame'>
```

```
RangeIndex: 7996 entries, 0 to 7995
```

```
Data columns (total 21 columns):
```

#	Column	Non-Null Count	Dtype
0	aluminium	7996 non-null	float64
1	ammonia	7996 non-null	float64
2	arsenic	7996 non-null	float64
3	barium	7996 non-null	float64
4	cadmium	7996 non-null	float64
5	chloramine	7996 non-null	float64
6	chromium	7996 non-null	float64
7	copper	7996 non-null	float64
8	flouride	7996 non-null	float64
9	bacteria	7996 non-null	float64
10	viruses	7996 non-null	float64
11	lead	7996 non-null	float64
12	nitrates	7996 non-null	float64
13	nitrites	7996 non-null	float64
14	mercury	7996 non-null	float64
15	perchlorate	7996 non-null	float64
16	radium	7996 non-null	float64
17	selenium	7996 non-null	float64
18	silver	7996 non-null	float64
19	uranium	7996 non-null	float64
20	is_safe	7996 non-null	int64

```
dtypes: float64(20), int64(1)
```

```
memory usage: 1.3 MB
```

```
# Unique values of each column
```

```
df.nunique()
```

aluminium	495
ammonia	2563
arsenic	107
barium	480
cadmium	23
chloramine	812
chromium	91
copper	201
flouride	151
bacteria	101
viruses	61

```

lead          200
nitrates      1803
nitrites      280
mercury       11
perchlorate   2999
radium        735
selenium      11
silver        51
uranium       10
is_safe       2
dtype: int64

```

```

# Checking for null values
df.isnull().sum()

```

```

aluminium      0
ammonia        0
arsenic        0
barium         0
cadmium        0
chloramine     0
chromium       0
copper         0
flouride       0
bacteria       0
viruses        0
lead          0
nitrates       0
nitrites       0
mercury        0
perchlorate    0
radium         0
selenium       0
silver         0
uranium        0
is_safe        0
dtype: int64

```

```

# Statistical summary of the dataset
df.describe(include='all')

```

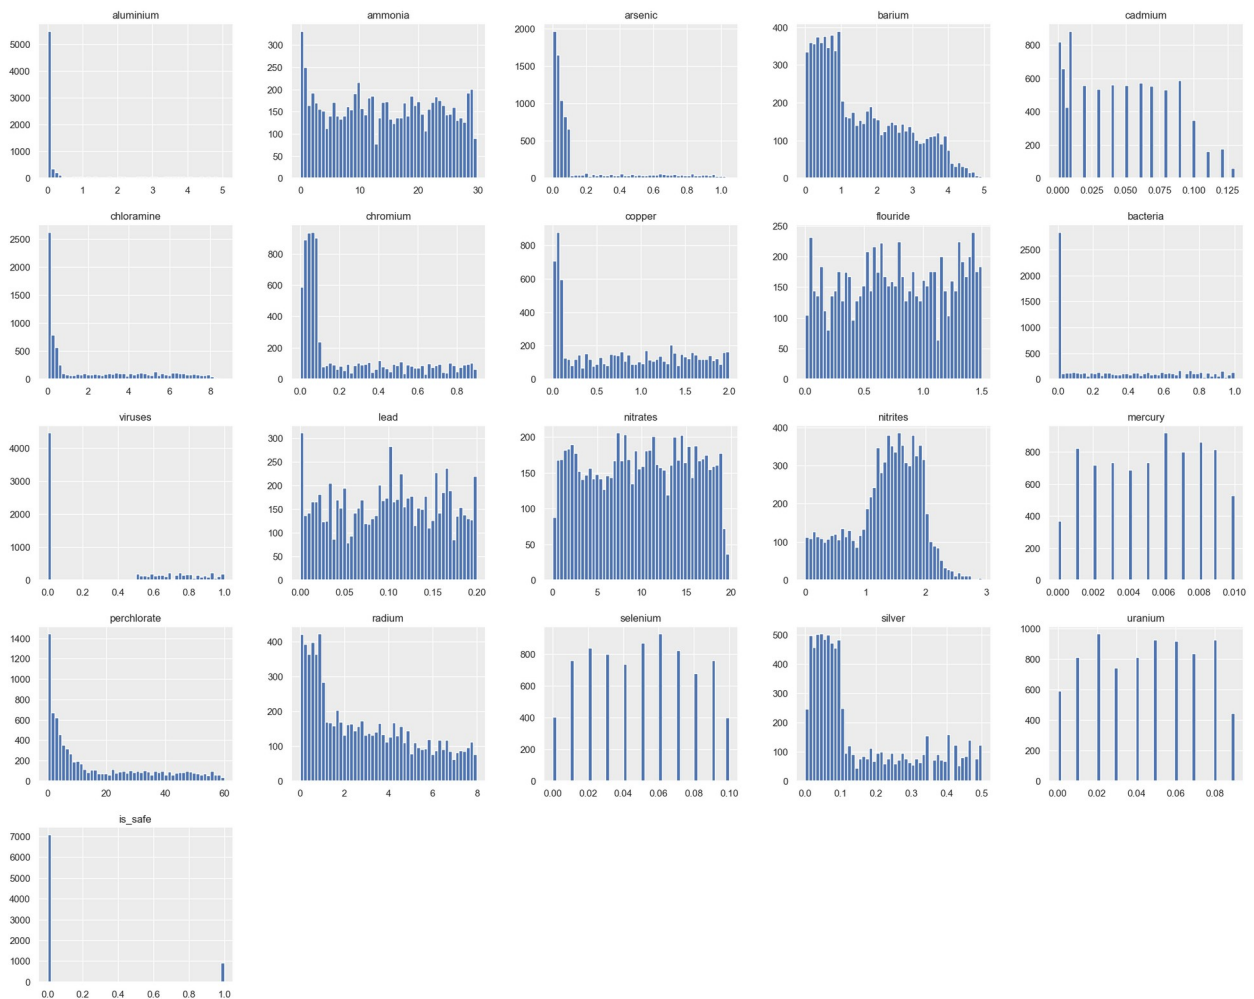
	aluminium	ammonia	arsenic	barium	cadmium
\count	7996.000000	7996.000000	7996.000000	7996.000000	7996.000000
mean	0.666396	14.278212	0.161477	1.567928	0.042803
std	1.265323	8.878930	0.252632	1.216227	0.036049
min	0.000000	-0.080000	0.000000	0.000000	0.000000

25%	0.040000	6.577500	0.030000	0.560000	0.008000
50%	0.070000	14.130000	0.050000	1.190000	0.040000
75%	0.280000	22.132500	0.100000	2.482500	0.070000
max	5.050000	29.840000	1.050000	4.940000	0.130000
	chloramine	chromium	copper	flouride	bacteria
... \					
count	7996.000000	7996.000000	7996.000000	7996.000000	7996.000000
...					
mean	2.177589	0.247300	0.805940	0.771646	0.319714
...					
std	2.567210	0.270663	0.653595	0.435423	0.329497
...					
min	0.000000	0.000000	0.000000	0.000000	0.000000
...					
25%	0.100000	0.050000	0.090000	0.407500	0.000000
...					
50%	0.530000	0.090000	0.750000	0.770000	0.220000
...					
75%	4.240000	0.440000	1.390000	1.160000	0.610000
...					
max	8.680000	0.900000	2.000000	1.500000	1.000000
...					
	lead	nitrates	nitrites	mercury	perchlorate
\					
count	7996.000000	7996.000000	7996.000000	7996.000000	7996.000000
mean	0.099431	9.819250	1.329846	0.005193	16.465266
std	0.058169	5.541977	0.573271	0.002967	17.688827
min	0.000000	0.000000	0.000000	0.000000	0.000000
25%	0.048000	5.000000	1.000000	0.003000	2.170000
50%	0.102000	9.930000	1.420000	0.005000	7.745000
75%	0.151000	14.610000	1.760000	0.008000	29.487500
max	0.200000	19.830000	2.930000	0.010000	60.010000
	radium	selenium	silver	uranium	is_safe
count	7996.000000	7996.000000	7996.000000	7996.000000	7996.000000

mean	2.920106	0.049684	0.147811	0.044672	0.114057
std	2.322805	0.028773	0.143569	0.026906	0.317900
min	0.000000	0.000000	0.000000	0.000000	0.000000
25%	0.820000	0.020000	0.040000	0.020000	0.000000
50%	2.410000	0.050000	0.080000	0.050000	0.000000
75%	4.670000	0.070000	0.240000	0.070000	0.000000
max	7.990000	0.100000	0.500000	0.090000	1.000000
[8 rows x 21 columns]					

Visualization

```
df.hist(bins=50, figsize=(25,20));
```

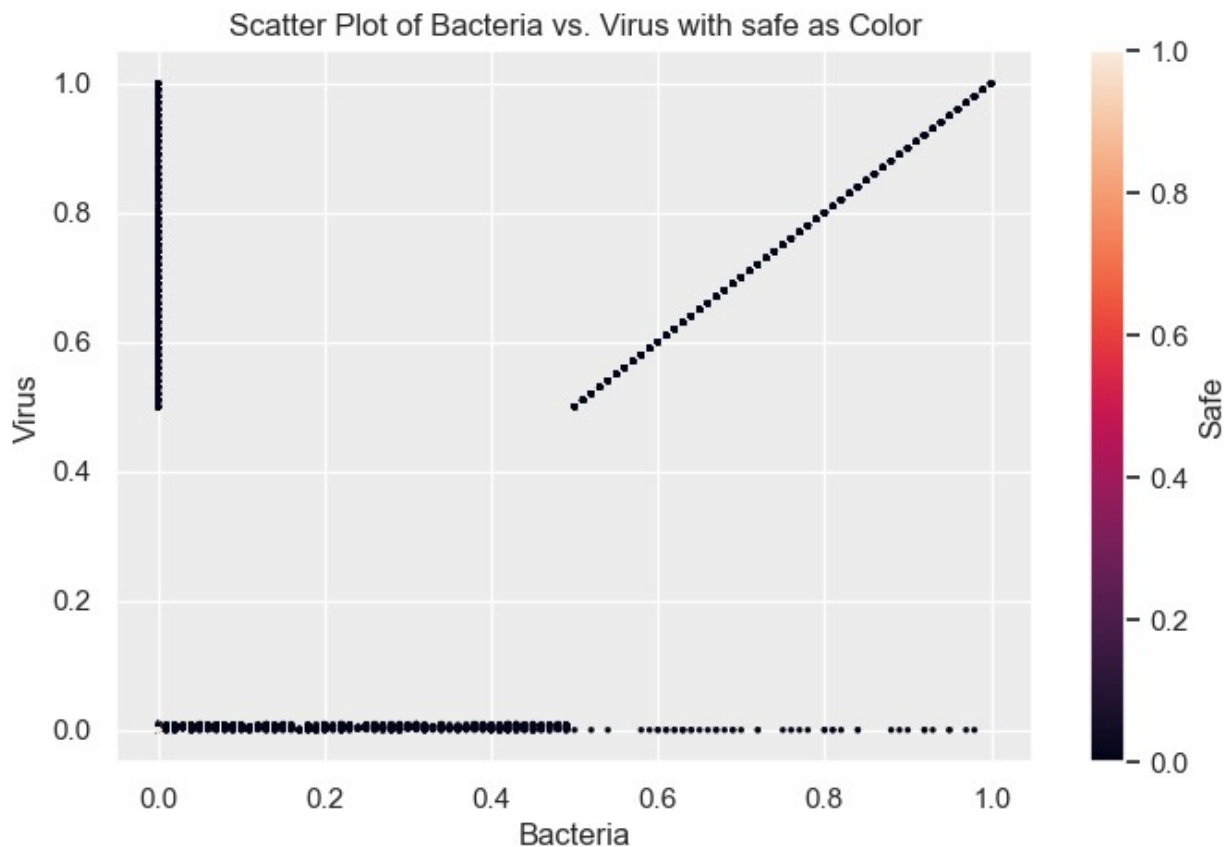


```

# Plotting a scatter plot for longitude and latitude with the house
prices as color
X=np.array(df)
bacteria=X[:,9]
virus=X[:,10]
is_safe=X[:,-1]

plt.figure(figsize=(8,5))
scatter = plt.scatter(bacteria,virus, c=is_safe, s=2)
plt.colorbar(scatter, label='Safe')
plt.xlabel('Bacteria')
plt.ylabel('Virus')
plt.title('Scatter Plot of Bacteria vs. Virus with safe as Color')
plt.show()

```



```

# Identify the target variable
target = 'is_safe'

# Separate the independent variables
independent_vars = df.columns.difference([target])

palette = sns.color_palette("husl", len(independent_vars))

```




```
# Create box plots for all independent variables with different colors
for i, var in enumerate(independent_vars):
    plt.figure(figsize=(10, 6))
    sns.boxplot(x=target, y=var, data=df, palette=[palette[i]])
    plt.title(f'Box Plot of {var} vs {target}')
    plt.xlabel(target)
    plt.ylabel(var)
    plt.show();
```

C:\Users\maadi\AppData\Local\Temp\ipykernel_19972\1591437457.py:12:
FutureWarning:

Passing `palette` without assigning `hue` is deprecated and will be removed in v0.14.0. Assign the `x` variable to `hue` and set `legend=False` for the same effect.

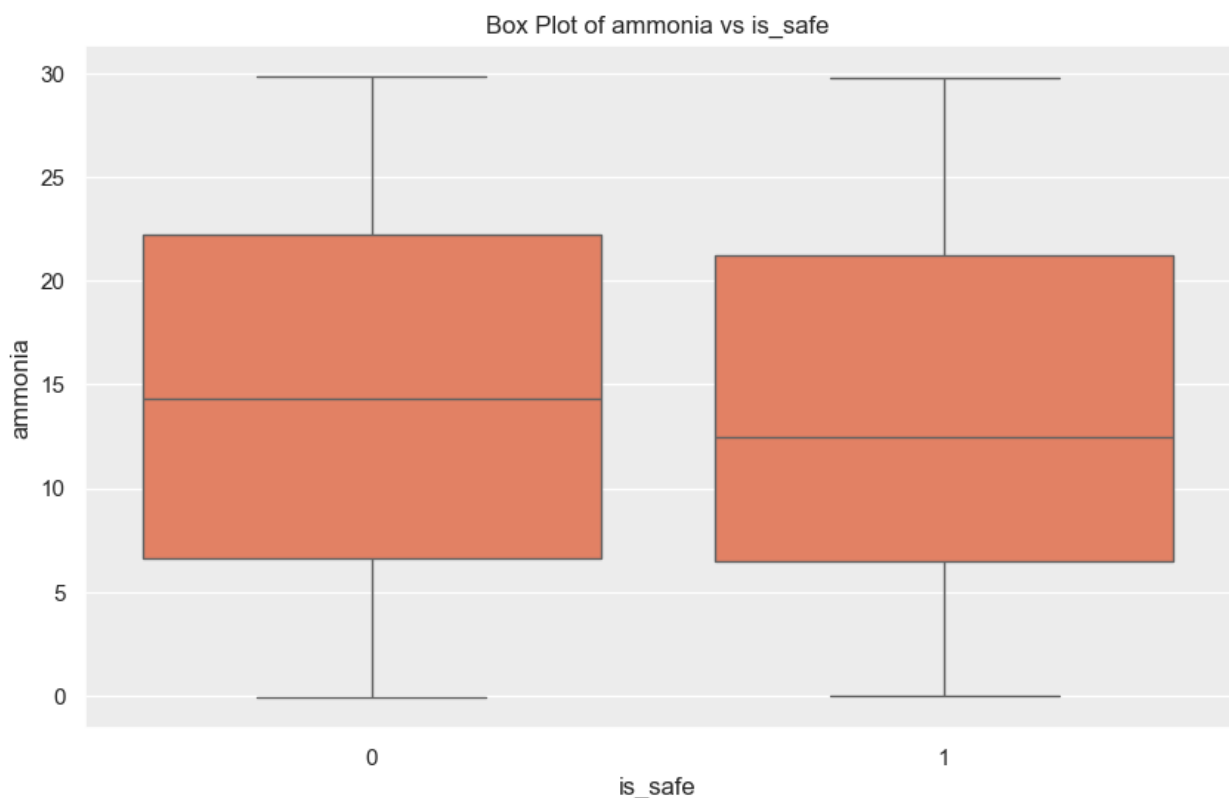
```
sns.boxplot(x=target, y=var, data=df, palette=[palette[i]])
C:\Users\maadi\AppData\Local\Temp\ipykernel_19972\1591437457.py:12:  
UserWarning:  
The palette list has fewer values (1) than needed (2) and will cycle,  
which may produce an uninterpretable plot.  
sns.boxplot(x=target, y=var, data=df, palette=[palette[i]])
```



```
C:\Users\maadi\AppData\Local\Temp\ipykernel_19972\1591437457.py:12:
FutureWarning:
```

Passing `palette` without assigning `hue` is deprecated and will be removed in v0.14.0. Assign the `x` variable to `hue` and set `legend=False` for the same effect.

```
sns.boxplot(x=target, y=var, data=df, palette=[palette[i]])
C:\Users\maadi\AppData\Local\Temp\ipykernel_19972\1591437457.py:12:
UserWarning:
The palette list has fewer values (1) than needed (2) and will cycle,
which may produce an uninterpretable plot.
sns.boxplot(x=target, y=var, data=df, palette=[palette[i]])
```




```
C:\Users\maadi\AppData\Local\Temp\ipykernel_19972\1591437457.py:12:
FutureWarning:
```

Passing `palette` without assigning `hue` is deprecated and will be removed in v0.14.0. Assign the `x` variable to `hue` and set `legend=False` for the same effect.

```
sns.boxplot(x=target, y=var, data=df, palette=[palette[i]])
C:\Users\maadi\AppData\Local\Temp\ipykernel_19972\1591437457.py:12:
UserWarning:
The palette list has fewer values (1) than needed (2) and will cycle,
```

which may produce an uninterpretable plot.

```
sns.boxplot(x=target, y=var, data=df, palette=[palette[i]])
```



C:\Users\maadi\AppData\Local\Temp\ipykernel_19972\1591437457.py:12:
FutureWarning:


Passing `palette` without assigning `hue` is deprecated and will be removed in v0.14.0. Assign the `x` variable to `hue` and set `legend=False` for the same effect.

```
sns.boxplot(x=target, y=var, data=df, palette=[palette[i]])
```

C:\Users\maadi\AppData\Local\Temp\ipykernel_19972\1591437457.py:12:
UserWarning:

The palette list has fewer values (1) than needed (2) and will cycle, which may produce an uninterpretable plot.

```
sns.boxplot(x=target, y=var, data=df, palette=[palette[i]])
```



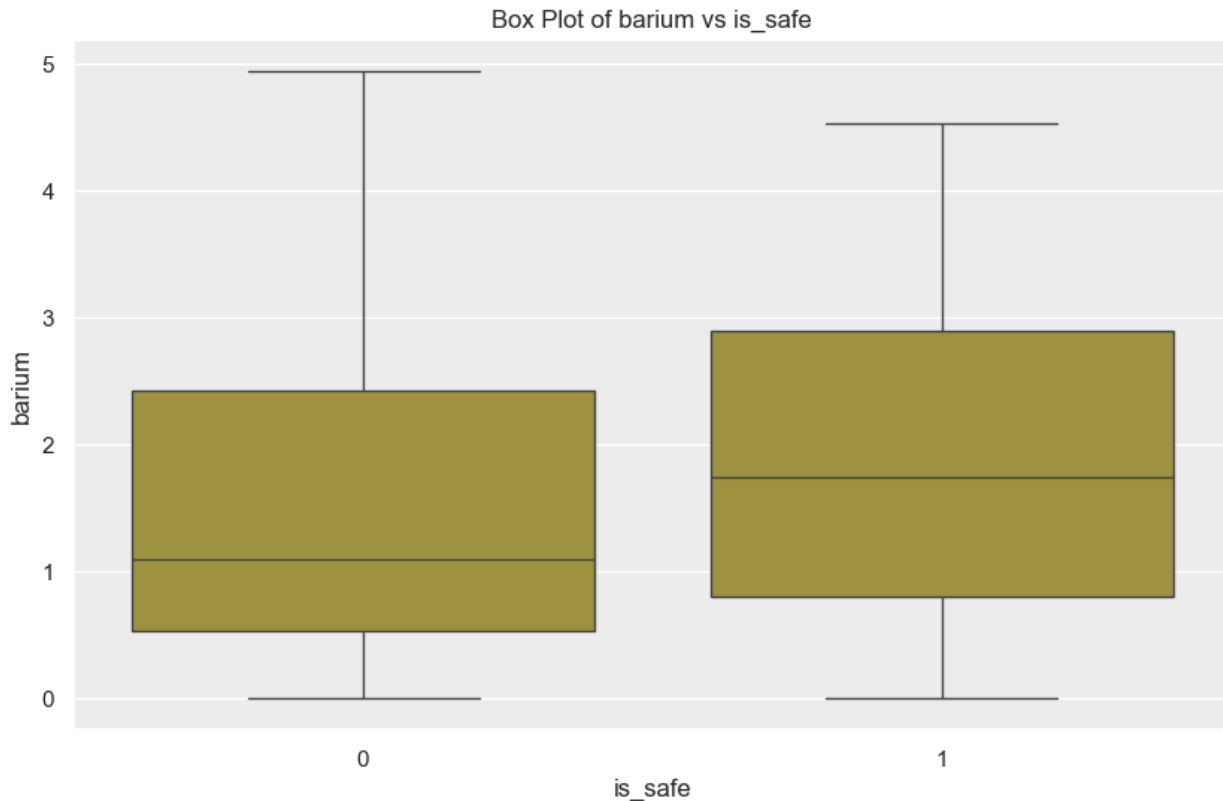
```
C:\Users\maadi\AppData\Local\Temp\ipykernel_19972\1591437457.py:12:
FutureWarning:
```

Passing `palette` without assigning `hue` is deprecated and will be removed in v0.14.0. Assign the `x` variable to `hue` and set `legend=False` for the same effect.

```
sns.boxplot(x=target, y=var, data=df, palette=[palette[i]])
C:\Users\maadi\AppData\Local\Temp\ipykernel_19972\1591437457.py:12:
UserWarning:
```

The palette list has fewer values (1) than needed (2) and will cycle, which may produce an uninterpretable plot.

```
sns.boxplot(x=target, y=var, data=df, palette=[palette[i]])
```




```
C:\Users\maadi\AppData\Local\Temp\ipykernel_19972\1591437457.py:12:
FutureWarning:
```

```
Passing `palette` without assigning `hue` is deprecated and will be
removed in v0.14.0. Assign the `x` variable to `hue` and set
`legend=False` for the same effect.
```

```
sns.boxplot(x=target, y=var, data=df, palette=[palette[i]])
C:\Users\maadi\AppData\Local\Temp\ipykernel_19972\1591437457.py:12:
UserWarning:
```

```
The palette list has fewer values (1) than needed (2) and will cycle,
which may produce an uninterpretable plot.
```

```
sns.boxplot(x=target, y=var, data=df, palette=[palette[i]])
```



```
C:\Users\maadi\AppData\Local\Temp\ipykernel_19972\1591437457.py:12:
FutureWarning:
```


```
Passing `palette` without assigning `hue` is deprecated and will be
removed in v0.14.0. Assign the `x` variable to `hue` and set
`legend=False` for the same effect.
```

```
sns.boxplot(x=target, y=var, data=df, palette=[palette[i]])
```

```
C:\Users\maadi\AppData\Local\Temp\ipykernel_19972\1591437457.py:12:
UserWarning:
```

```
The palette list has fewer values (1) than needed (2) and will cycle,
which may produce an uninterpretable plot.
```

```
sns.boxplot(x=target, y=var, data=df, palette=[palette[i]])
```




```
C:\Users\maadi\AppData\Local\Temp\ipykernel_19972\1591437457.py:12:
FutureWarning:
```

```
Passing `palette` without assigning `hue` is deprecated and will be
removed in v0.14.0. Assign the `x` variable to `hue` and set
`legend=False` for the same effect.
```

```
sns.boxplot(x=target, y=var, data=df, palette=[palette[i]])
C:\Users\maadi\AppData\Local\Temp\ipykernel_19972\1591437457.py:12:
UserWarning:
```

```
The palette list has fewer values (1) than needed (2) and will cycle,
which may produce an uninterpretable plot.
```

```
sns.boxplot(x=target, y=var, data=df, palette=[palette[i]])
```



```
C:\Users\maadi\AppData\Local\Temp\ipykernel_19972\1591437457.py:12:
FutureWarning:
```

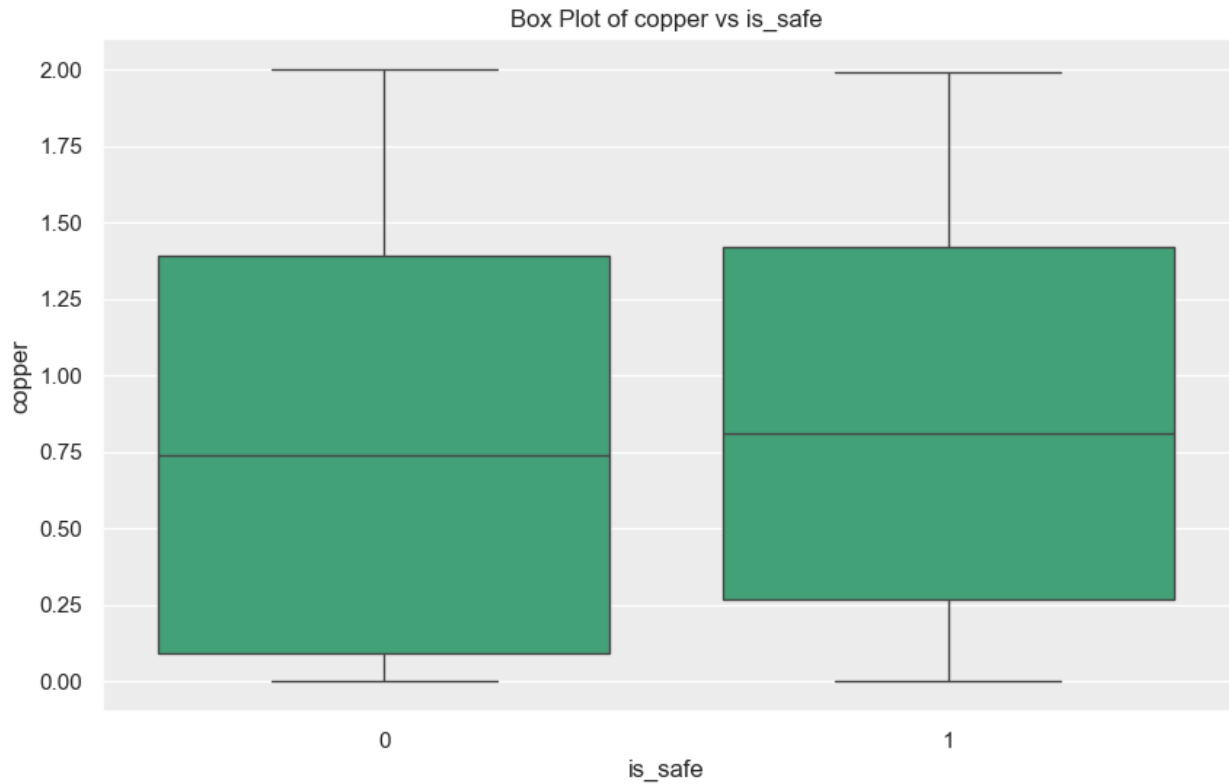
Passing `palette` without assigning `hue` is deprecated and will be removed in v0.14.0. Assign the `x` variable to `hue` and set `legend=False` for the same effect.

```
sns.boxplot(x=target, y=var, data=df, palette=[palette[i]])
```

```
C:\Users\maadi\AppData\Local\Temp\ipykernel_19972\1591437457.py:12:
UserWarning:
```

The palette list has fewer values (1) than needed (2) and will cycle, which may produce an uninterpretable plot.

```
sns.boxplot(x=target, y=var, data=df, palette=[palette[i]])
```

```
C:\Users\maadi\AppData\Local\Temp\ipykernel_19972\1591437457.py:12:
FutureWarning:
```


```
Passing `palette` without assigning `hue` is deprecated and will be
removed in v0.14.0. Assign the `x` variable to `hue` and set
`legend=False` for the same effect.
```

```
sns.boxplot(x=target, y=var, data=df, palette=[palette[i]])
```

```
C:\Users\maadi\AppData\Local\Temp\ipykernel_19972\1591437457.py:12:
UserWarning:
```

```
The palette list has fewer values (1) than needed (2) and will cycle,
which may produce an uninterpretable plot.
```

```
sns.boxplot(x=target, y=var, data=df, palette=[palette[i]])
```



```
C:\Users\maadi\AppData\Local\Temp\ipykernel_19972\1591437457.py:12:
FutureWarning:
```


Passing `palette` without assigning `hue` is deprecated and will be removed in v0.14.0. Assign the `x` variable to `hue` and set `legend=False` for the same effect.

```
sns.boxplot(x=target, y=var, data=df, palette=[palette[i]])
```

```
C:\Users\maadi\AppData\Local\Temp\ipykernel_19972\1591437457.py:12:
UserWarning:
```

The palette list has fewer values (1) than needed (2) and will cycle, which may produce an uninterpretable plot.

```
sns.boxplot(x=target, y=var, data=df, palette=[palette[i]])
```




```
C:\Users\maadi\AppData\Local\Temp\ipykernel_19972\1591437457.py:12:
FutureWarning:
```

```
Passing `palette` without assigning `hue` is deprecated and will be
removed in v0.14.0. Assign the `x` variable to `hue` and set
`legend=False` for the same effect.
```

```
sns.boxplot(x=target, y=var, data=df, palette=[palette[i]])
C:\Users\maadi\AppData\Local\Temp\ipykernel_19972\1591437457.py:12:
UserWarning:
The palette list has fewer values (1) than needed (2) and will cycle,
which may produce an uninterpretable plot.
```


```
sns.boxplot(x=target, y=var, data=df, palette=[palette[i]])
```



```
C:\Users\maadi\AppData\Local\Temp\ipykernel_19972\1591437457.py:12:
FutureWarning:
```

```
Passing `palette` without assigning `hue` is deprecated and will be
removed in v0.14.0. Assign the `x` variable to `hue` and set
`legend=False` for the same effect.
```

```
sns.boxplot(x=target, y=var, data=df, palette=[palette[i]])
C:\Users\maadi\AppData\Local\Temp\ipykernel_19972\1591437457.py:12:
UserWarning:
The palette list has fewer values (1) than needed (2) and will cycle,
which may produce an uninterpretable plot.
sns.boxplot(x=target, y=var, data=df, palette=[palette[i]])
```



```
C:\Users\maadi\AppData\Local\Temp\ipykernel_19972\1591437457.py:12:
FutureWarning:
```


```
Passing `palette` without assigning `hue` is deprecated and will be
removed in v0.14.0. Assign the `x` variable to `hue` and set
`legend=False` for the same effect.
```

```
sns.boxplot(x=target, y=var, data=df, palette=[palette[i]])
```

```
C:\Users\maadi\AppData\Local\Temp\ipykernel_19972\1591437457.py:12:
UserWarning:
```

```
The palette list has fewer values (1) than needed (2) and will cycle,
which may produce an uninterpretable plot.
```

```
sns.boxplot(x=target, y=var, data=df, palette=[palette[i]])
```



```
C:\Users\maadi\AppData\Local\Temp\ipykernel_19972\1591437457.py:12:
FutureWarning:
```

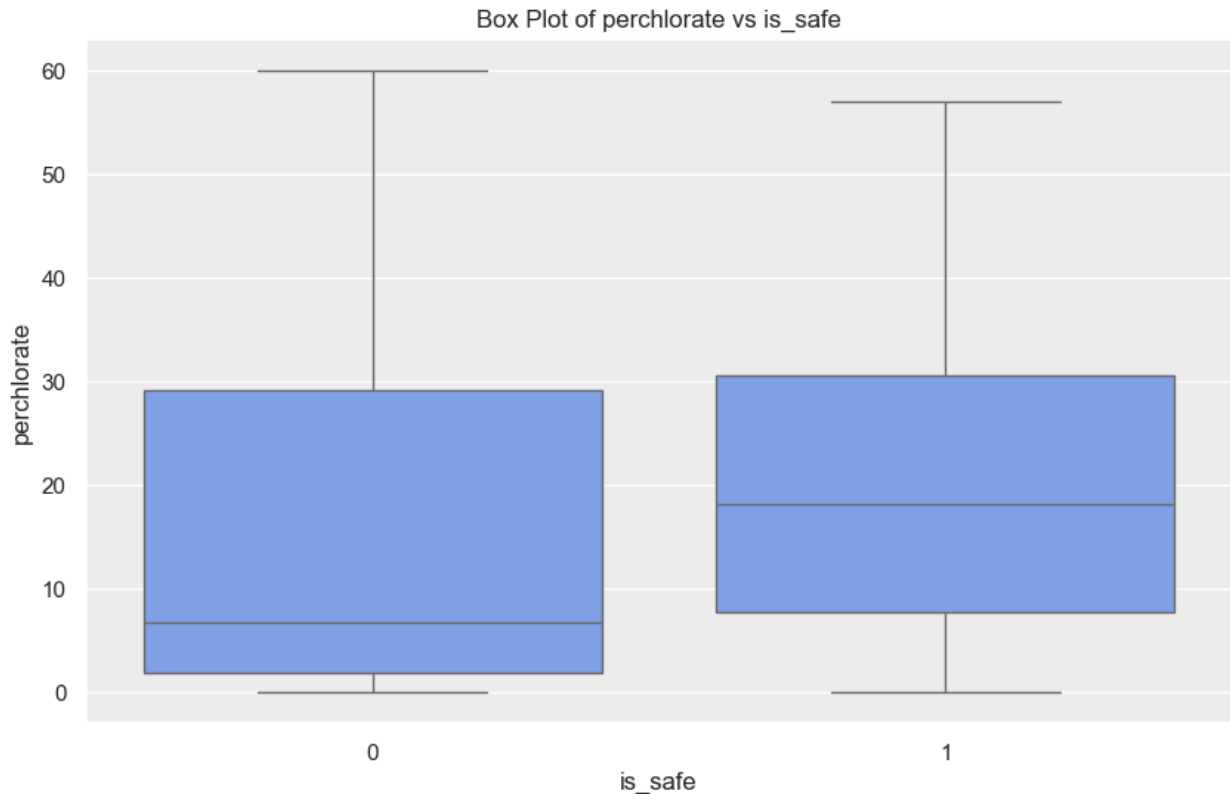
Passing `palette` without assigning `hue` is deprecated and will be removed in v0.14.0. Assign the `x` variable to `hue` and set `legend=False` for the same effect.

```
sns.boxplot(x=target, y=var, data=df, palette=[palette[i]])
```

```
C:\Users\maadi\AppData\Local\Temp\ipykernel_19972\1591437457.py:12:
UserWarning:
```

The palette list has fewer values (1) than needed (2) and will cycle, which may produce an uninterpretable plot.

```
sns.boxplot(x=target, y=var, data=df, palette=[palette[i]])
```



```
C:\Users\maadi\AppData\Local\Temp\ipykernel_19972\1591437457.py:12:
FutureWarning:
```


Passing `palette` without assigning `hue` is deprecated and will be removed in v0.14.0. Assign the `x` variable to `hue` and set `legend=False` for the same effect.

```
sns.boxplot(x=target, y=var, data=df, palette=[palette[i]])
```

```
C:\Users\maadi\AppData\Local\Temp\ipykernel_19972\1591437457.py:12:
UserWarning:
```

The palette list has fewer values (1) than needed (2) and will cycle, which may produce an uninterpretable plot.

```
sns.boxplot(x=target, y=var, data=df, palette=[palette[i]])
```



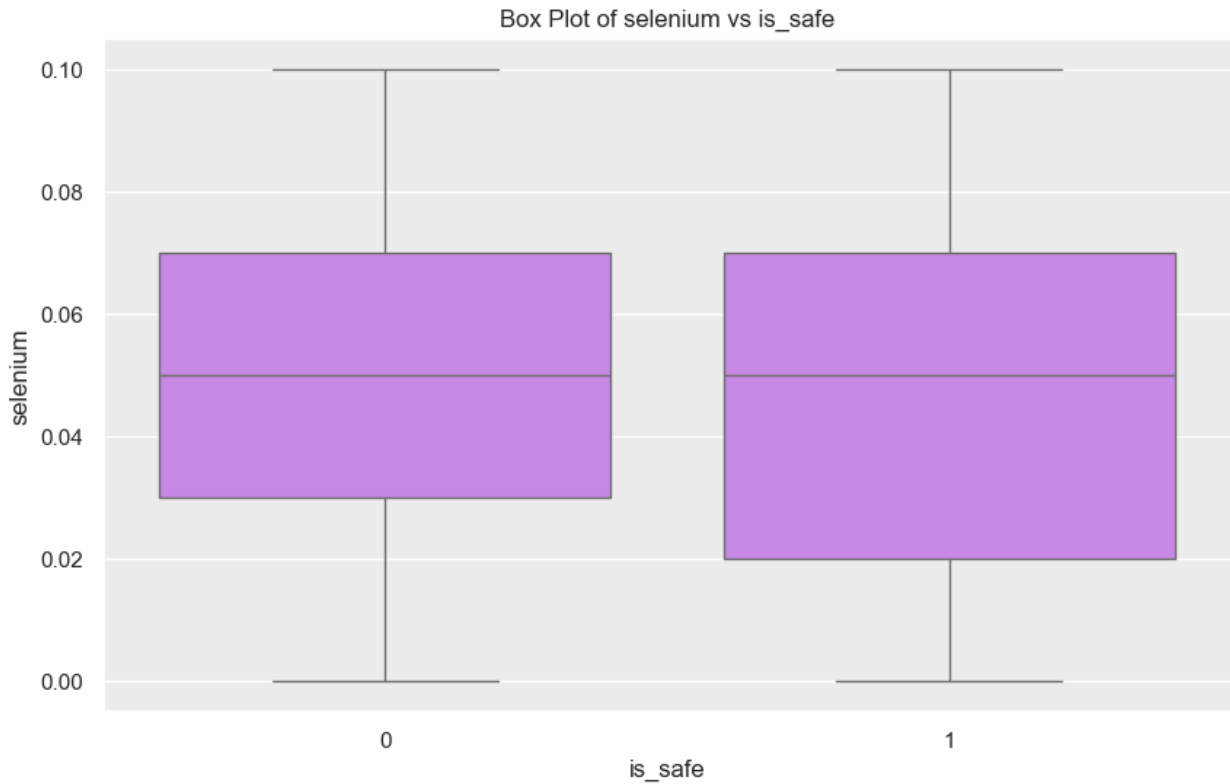
```
C:\Users\maadi\AppData\Local\Temp\ipykernel_19972\1591437457.py:12:
FutureWarning:
```

Passing `palette` without assigning `hue` is deprecated and will be removed in v0.14.0. Assign the `x` variable to `hue` and set `legend=False` for the same effect.

```
sns.boxplot(x=target, y=var, data=df, palette=[palette[i]])
C:\Users\maadi\AppData\Local\Temp\ipykernel_19972\1591437457.py:12:
UserWarning:
```

The palette list has fewer values (1) than needed (2) and will cycle, which may produce an uninterpretable plot.

```
sns.boxplot(x=target, y=var, data=df, palette=[palette[i]])
```

```
C:\Users\maadi\AppData\Local\Temp\ipykernel_19972\1591437457.py:12:
FutureWarning:
```


```
Passing `palette` without assigning `hue` is deprecated and will be
removed in v0.14.0. Assign the `x` variable to `hue` and set
`legend=False` for the same effect.
```

```
sns.boxplot(x=target, y=var, data=df, palette=[palette[i]])
```

```
C:\Users\maadi\AppData\Local\Temp\ipykernel_19972\1591437457.py:12:
UserWarning:
```

```
The palette list has fewer values (1) than needed (2) and will cycle,
which may produce an uninterpretable plot.
```

```
sns.boxplot(x=target, y=var, data=df, palette=[palette[i]])
```




```
C:\Users\maadi\AppData\Local\Temp\ipykernel_19972\1591437457.py:12:
FutureWarning:
```

Passing `palette` without assigning `hue` is deprecated and will be removed in v0.14.0. Assign the `x` variable to `hue` and set `legend=False` for the same effect.

```
sns.boxplot(x=target, y=var, data=df, palette=[palette[i]])
C:\Users\maadi\AppData\Local\Temp\ipykernel_19972\1591437457.py:12:
UserWarning:
```

The palette list has fewer values (1) than needed (2) and will cycle, which may produce an uninterpretable plot.

```
sns.boxplot(x=target, y=var, data=df, palette=[palette[i]])
```



```
C:\Users\maadi\AppData\Local\Temp\ipykernel_19972\1591437457.py:12:
FutureWarning:
```

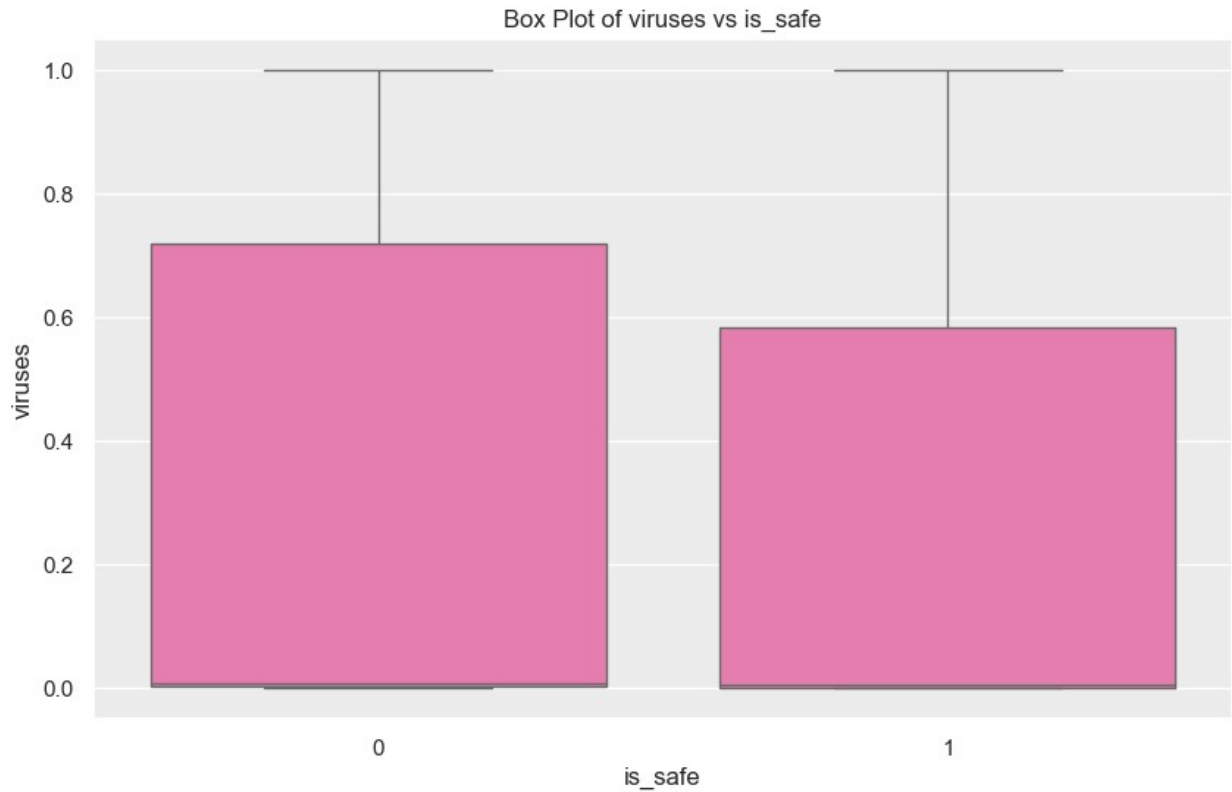
```
Passing `palette` without assigning `hue` is deprecated and will be
removed in v0.14.0. Assign the `x` variable to `hue` and set
`legend=False` for the same effect.
```

```
sns.boxplot(x=target, y=var, data=df, palette=[palette[i]])
```


```
C:\Users\maadi\AppData\Local\Temp\ipykernel_19972\1591437457.py:12:
UserWarning:
```

```
The palette list has fewer values (1) than needed (2) and will cycle,
which may produce an uninterpretable plot.
```

```
sns.boxplot(x=target, y=var, data=df, palette=[palette[i]])
```

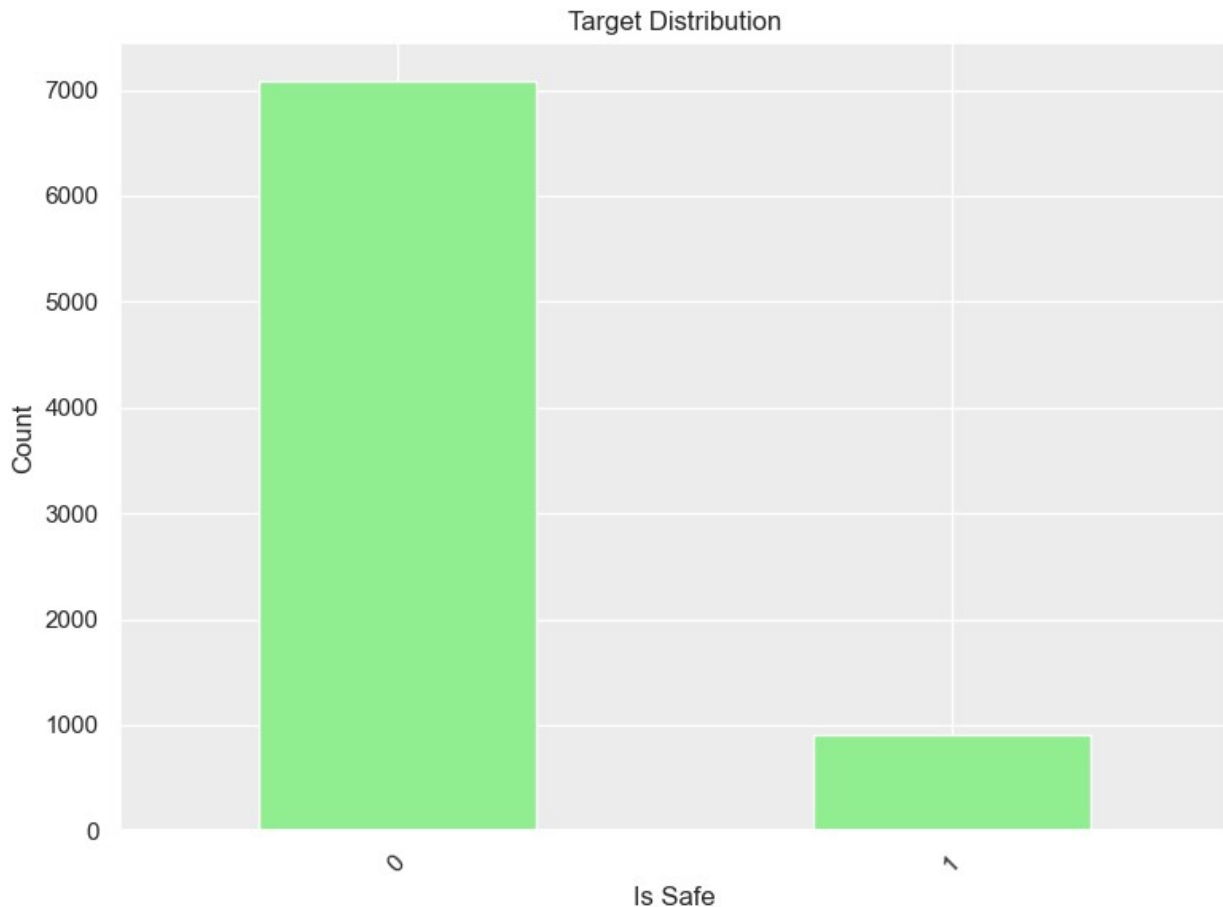


```
# Plotting heatmap for correlation
df_corr = df.corr()
plt.figure(figsize = (15, 10))
sns.heatmap(df_corr, fmt = ".2f", annot = True, cmap = "YlGnBu")
plt.show()
```



```
is_safe = df['is_safe'].value_counts()

# Plotting the bar plot
plt.figure(figsize=(8, 6))
is_safe.plot(kind='bar', color='lightgreen')
plt.title('Target Distribution')
plt.xlabel('Is Safe')
plt.ylabel('Count')
plt.xticks(rotation=45)
plt.tight_layout()
plt.show()
```



Model Building

```
X=np.array(df)
features=X[:,0:19]
target=X[:, -1]

# Splitting data for model building
X_train, X_test, y_train, y_test = train_test_split(features, target,
test_size=0.2, random_state=42)

# Fitting the model
model = KNeighborsClassifier()
model.fit(X_train,y_train)

# Predicting for training set
train_pred = model.predict(X_train)

# Calculating training accuracy
train_accuracy = accuracy_score(y_train,train_pred)
print("Training accuracy is: " + str(train_accuracy))

# Getting confusion matrix and Classification report
print("Confusion Matrix: ")
```

```

print(confusion_matrix(y_train, train_pred))

print("Classification Report: ")
print(classification_report(y_train, train_pred))

Training accuracy is: 0.9111944965603502
Confusion Matrix:
[[5620   64]
 [ 504  208]]
Classification Report:

```

	precision	recall	f1-score	support
0.0	0.92	0.99	0.95	5684
1.0	0.76	0.29	0.42	712
accuracy			0.91	6396
macro avg	0.84	0.64	0.69	6396
weighted avg	0.90	0.91	0.89	6396

```


# Adding jitter to avoid overlap in scatter plot
def add_jitter(arr, jitter_strength=0.1):
    stdev = jitter_strength * (max(arr) - min(arr))
    return arr + np.random.randn(len(arr)) * stdev

# Plotting the scatter plot for actual vs predicted values
plt.figure(figsize=(10, 7))
sns.scatterplot(x=add_jitter(y_train), y=add_jitter(train_pred),
marker='o', alpha=0.6)

# Adding a diagonal line to show perfect prediction
plt.plot([min(y_train), max(y_train)], [min(y_train), max(y_train)],
color='red', linestyle='--')

plt.xlabel('Actual Values')
plt.ylabel('Predicted Values')
plt.title('Scatter Plot of Actual vs. Predicted Values with Jitter')
plt.show()

```



```
# Predicting for testing set
test_pred = model.predict(X_test)

# Calculating testing accuracy
test_accuracy = accuracy_score(y_test, test_pred)
print("Training accuracy is: " + str(test_accuracy))

# Getting confusion matrix and Classification report
print("Confusion Matrix: ")
print(confusion_matrix(y_test, test_pred))

print("Classification Report: ")
print(classification_report(y_test, test_pred))
```

Training accuracy is: 0.8625

Confusion Matrix:

```
[[1359  41]
 [ 179  21]]
```

Classification Report:

	precision	recall	f1-score	support
0.0	0.88	0.97	0.93	1400
1.0	0.34	0.10	0.16	200

accuracy			0.86	1600
macro avg	0.61	0.54	0.54	1600
weighted avg	0.82	0.86	0.83	1600

```
# Plotting the scatter plot for actual vs predicted values
```

```
plt.figure(figsize=(10, 7))
sns.scatterplot(x=add_jitter(y_test), y=add_jitter(test_pred),
marker='o', alpha=0.6)
```

```
# Adding a diagonal line to show perfect prediction
```


```
plt.plot([min(y_test), max(y_test)], [min(y_test), max(y_test)],
color='red', linestyle='--')
```

```
plt.xlabel('Actual Values')
```

```
plt.ylabel('Predicted Values')
```

```
plt.title('Scatter Plot of Actual vs. Predicted Values with Jitter')
```

```
plt.show()
```



Model Improvement

```
n_components = 12

# Importing PCA
from sklearn.decomposition import PCA

# Let's say, components = 12
pca = PCA(n_components=12)
pca.fit(X)
x_pca = pca.transform(X)

# Create the dataframe
df_pca1 = pd.DataFrame(x_pca, columns=['PC{}'.format(i+1) for i in
range(n_components)])
print(df_pca1)
```

	PC1	PC2	PC3	PC4	PC5	PC6
PC7 \						
0	20.875094	6.460591	6.377741	-0.075213	-4.468788	1.176144
0.683676						
1	16.518599	-5.860612	-7.776526	1.008795	1.832076	1.220882
0.589493						
2	33.872073	2.291372	4.502261	0.868194	-2.580851	-1.195053
1.749896						
3	-6.982723	2.501815	-8.403896	3.420961	5.031750	0.321890
0.619224						
4	1.035591	-9.979904	-3.154869	-0.397186	0.542414	-0.216626
1.444306						
...
...						
7991	-13.441977	5.673185	4.448488	-0.555633	-0.545002	0.212029
0.908805						
7992	-14.563406	-10.827515	0.293340	-1.299944	0.241330	-0.217025
0.346170						
7993	-15.619716	6.448067	6.054763	1.461705	-2.448939	-0.271492
0.401118						
7994	-16.733170	3.333144	-9.856313	-0.688998	2.457217	-0.059811
0.824874						
7995	-15.618424	6.448248	6.054627	1.466486	-2.445993	-0.290451
0.302166						
	PC8	PC9	PC10	PC11	PC12	
0	-0.786923	-0.261715	-0.484511	0.542568	0.956504	
1	-0.233650	0.323953	0.420361	0.016980	0.811942	
2	-0.955396	-0.284386	-0.207000	-0.329719	-0.045636	
3	0.624846	-0.557999	0.482380	-0.032048	0.675337	
4	-0.354707	0.007436	-0.437549	-0.049879	0.839236	
...	
7991	-0.772219	-0.173152	-0.049259	-0.699844	0.889071	
7992	-0.757892	0.083607	-0.073058	-0.863743	0.860983	

```

7993 -0.756346 -0.107219 -0.241161 -0.262259 0.881928
7994 -1.028636 -0.996383 -0.697608 0.583452 0.748330
7995 -0.742315 -0.124202 -0.206084 -0.344706 0.879444

```

```
[7996 rows x 12 columns]
```

```

explained_variance = pca.explained_variance_
explained_variance_ratio = pca.explained_variance_ratio_
cumulative_explained_variance = np.cumsum(explained_variance_ratio)

```

```

# Plot the explained variance ratio and cumulative explained variance
plt.figure(figsize=(10, 5))

```

```
# Plot explained variance ratio
```

```

plt.subplot(1, 2, 1)
plt.bar(range(1, len(explained_variance_ratio) + 1),
        explained_variance_ratio, alpha=0.5, align='center')
plt.xlabel('Principal Component')
plt.ylabel('Explained Variance Ratio')
plt.title('Explained Variance Ratio')

```

```
# Plot cumulative explained variance
```

```


plt.subplot(1, 2, 2)
plt.plot(range(1, len(cumulative_explained_variance) + 1),
        cumulative_explained_variance, marker='o', linestyle='-')
plt.xlabel('Number of Principal Components')
plt.ylabel('Cumulative Explained Variance')
plt.title('Cumulative Explained Variance')

```

```

plt.tight_layout()
plt.show()

```



```

target=X[:, -1]

# Splitting data for model building
im_X_train, im_X_test, im_y_train, im_y_test =
train_test_split(df_pcal, target, test_size=0.2, random_state=42)

# Fitting the model
model = KNeighborsClassifier()
model.fit(im_X_train, im_y_train)

# Predicting for training set
im_train_pred = model.predict(im_X_train)

# Calculating training accuracy
im_train_accuracy = accuracy_score(im_y_train, im_train_pred)
print("Training accuracy is: " + str(im_train_accuracy))

# Getting confusion matrix and Classification report
print("Confusion Matrix: ")
print(confusion_matrix(im_y_train, im_train_pred))

print("Classification Report: ")
print(classification_report(im_y_train, im_train_pred))

Training accuracy is: 0.9177611006879299
Confusion Matrix:
[[5632   52]
 [ 474  238]]
Classification Report:

```

	precision	recall	f1-score	support
0.0	0.92	0.99	0.96	5684
1.0	0.82	0.33	0.48	712
accuracy			0.92	6396
macro avg	0.87	0.66	0.72	6396
weighted avg	0.91	0.92	0.90	6396

```

# Predicting for testing set
im_test_pred = model.predict(im_X_test)

# Calculating testing accuracy
im_test_accuracy = accuracy_score(im_y_test, im_test_pred)
print("Training accuracy is: " + str(im_test_accuracy))

# Getting confusion matrix and Classification report
print("Confusion Matrix: ")
print(confusion_matrix(im_y_test, im_test_pred))

```

```

print("Classification Report: ")
print(classification_report(im_y_test,im_test_pred))

Training accuracy is: 0.8725
Confusion Matrix:
[[1369   31]
 [ 173   27]]
Classification Report:

```


	precision	recall	f1-score	support
0.0	0.89	0.98	0.93	1400
1.0	0.47	0.14	0.21	200
accuracy			0.87	1600
macro avg	0.68	0.56	0.57	1600
weighted avg	0.84	0.87	0.84	1600

```

comparisions = ['Without PCA' , 'With PCA']
accuracies = [test_accuracy, im_test_accuracy]

plt.figure(figsize=(10, 6))
plt.plot(comparisions, accuracies, marker='o', linestyle='--',
color='b', label='Testing Accuracy')
plt.ylim(0, 1)
plt.xlabel('Comparisions')
plt.ylabel('Accuracy')
plt.title('Comparison of Accuracy Scores for Testing Sets')
plt.legend()
for i, acc in enumerate(accuracies):
    plt.text(i, acc + 0.02, round(acc, 4), ha='center')
plt.show()

```



Hyperparameter Tuning

```
# create numpy array for future K value
neighbors = np.arange(1, 40)
train_accuracy = np.empty(len(neighbors))
test_accuracy = np.empty(len(neighbors))

# Code Help taken from :
https://www.kaggle.com/code/elisthefox/ultimate-guide-to-k-nearest-neighbors-k-nn#4.-Results-of-K-NN-implementation-after-preparation-of-data

# Loop over different values of k
for i, k in enumerate(neighbors):
    # Setup a k-NN Classifier with k neighbors: knn
    knn = KNeighborsClassifier(n_neighbors=k)

    # Fit the classifier to the training data
    knn.fit(im_X_train, im_y_train)

    # Compute accuracy on the training set
    train_accuracy[i] = knn.score(im_X_train, im_y_train)

    # Compute accuracy on the testing set
    test_accuracy[i] = knn.score(im_X_test, im_y_test)
```

```

# Generate plot
sns.set(rc={'axes.facecolor':'#ECECEC'}) #background color of all
plots
plt.figure(figsize=(12,6.5))
plt.title(label='K-NN: Varying Number of Neighbors', fontsize=15,
fontweight='bold', fontname='Verdana', ha='center')
plt.plot(neighbors, test_accuracy, label = 'Testing Accuracy',
color='#E68753')
plt.plot(neighbors, train_accuracy, label = 'Training Accuracy', color
= '#409996')
plt.legend()
plt.xlabel('Number of Neighbors')
plt.ylabel('Accuracy')
plt.show()

```

